

# Programming with Parallel Objects

## Session 1: Fundamentals of Parallel Objects Programming

Esteban Meneses

Advanced Computing Laboratory  
Costa Rica National High Technology Center

School of Computing  
Costa Rica Institute of Technology

2017

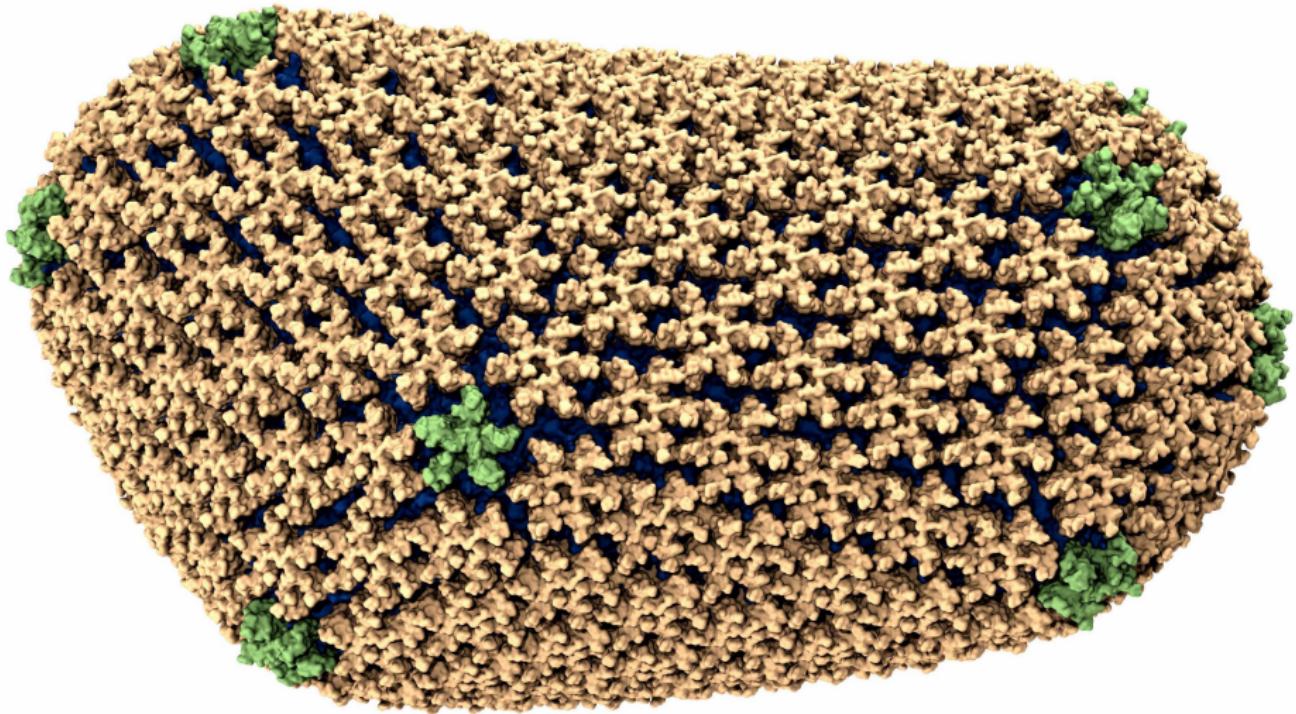


## **ACKNOWLEDGEMENTS**

Most of the content of this tutorial was provided courtesy of Prof. Laxmikant V. Kalé from the Parallel Programming Laboratory (PPL) of the University of Illinois at Urbana-Champaign. The PPL originally developed the ideas behind parallel object programming and has maintained the Charm++ code base for more than 20 years.

## **ADMINISTRATIVIA**

- ▶ Official Charm++ website:  
**<http://www.charmplusplus.org/>**
- ▶ Slides and code for this tutorial:  
**[https://github.com/emenesesrojas/parallel\\_objects.git](https://github.com/emenesesrojas/parallel_objects.git)**
- ▶ Questions on parallel objects programming:  
**<https://ecar2017.slack.com>**
- ▶ Advanced questions on Charm++:  
**ppl@cs.illinois.edu**
- ▶ Official instructor's email address:  
**emeneses@cenat.ac.cr**



## Structure of HIV capsid

Theoretical and Computational Biophysics Group  
University of Illinois at Urbana-Champaign

# Outline

Programming with  
Parallel Objects

Esteban Meneses

## Challenges in High Performance Computing

Introduction

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

Hello World

Benefits of Charm++

Charm++ Programming Language

Charm++ Basics

Object Collections

# Harnessing Parallelism

## Trends in system architecture

- ▶ Frequencies have stopped increasing
- ▶ Memory costs are high
  - ▶ Relatively low per core memory
- ▶ Increasing heterogeneity
  - ▶ Accelerators, simultaneous multithreading (SMT)
- ▶ Energy, power, and thermal considerations
- ▶ Frequent component failures

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

# Harnessing Parallelism

## High concurrency

- ▶ Compute resources are not faster cores, but **more cores**
- ▶ Unprecedented levels of available concurrency:

Rank*	Machine	Cores
1	NRCPC Sunway TaihuLight	10,649,600
2	NUDT Tianhe-2	3,120,000
5	IBM BG/Q Sequoia	1,572,864
8	Fujitsu K computer	705,024
6	Cray XC40 Cori	622,336

\* Top 500 list, June 2017

- ▶ Mid-size clusters will be ubiquitous
- ▶ Each thread of execution has to:
  - ▶ operate on lesser data
  - ▶ wait relatively longer for remote data
- ▶ Have to operate in **strong scaling** regime

# Harnessing Parallelism

## Next-generation applications

- ▶ Need for strong scaling
  - ▶ faster solutions (not just larger problems)
- ▶ Application characteristics
  - ▶ Multi-resolution
    - ▶ Adaptive, spatial and temporal resolutions
    - ▶ Dynamic/adaptive refinements: to handle application variation
  - ▶ Multi-module (multi-physics)
    - ▶ Complex physics in multiple, interacting modules
  - ▶ Adapt to a volatile computational environment
  - ▶ Exploit heterogeneous architecture
  - ▶ Deal with thermal and energy considerations
- ▶ Consequences:
  - ▶ Must support automated resource management
  - ▶ Must support interoperability and parallel composition

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion



# Harnessing Parallelism

## MPI programming model

- ▶ Highly successful
- ▶ Universally used
- ▶ Has supported application evolution from gigascale to petascale
- ▶ Library
- ▶ Communication primitives
- ▶ MPI does not directly support automated resource management (e.g. load balancing, fault tolerance, etc.)

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

# Stuff you Already Know

## Benefits of object-based code

- ▶ Objects encapsulate data
- ▶ Methods represent functionality relevant to that data
- ▶ Method invocations can modify/update state of the object/data
- ▶ Computation can be expressed in terms of objects interacting via method invocations
- ▶ Methods are natural units of sequential computation on object data
- ▶ Thoughtful design yields focused methods with single purpose
- ▶ Naturally expresses an object's response to inputs (signals/data)
- ▶ Nothing new
- ▶ Still quite uncommon in HPC code
- ▶ It's not about language syntax, it's about program structure

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

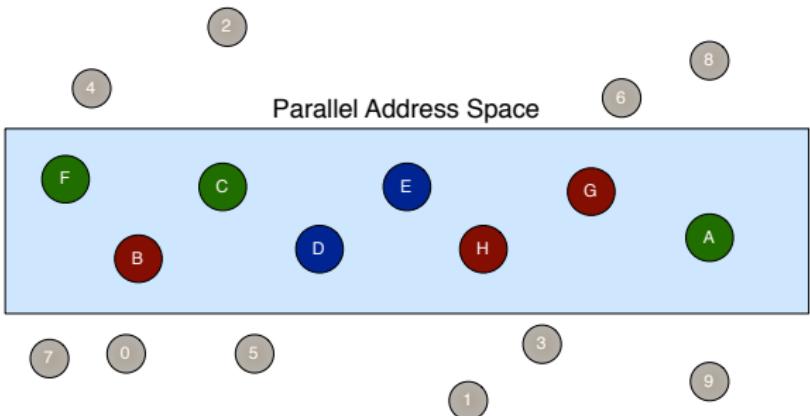
Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion



- ▶ Certain “special” object *instances* are:
  - ▶ first-class citizens in the parallel address space
  - ▶ with unique location-independent names
- ▶ Under the hood, the runtime handles locality and provides the mechanisms to promote objects to the parallel space

# Globally-visible Methods

## Entry methods

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design

Execution Mode

Hello World

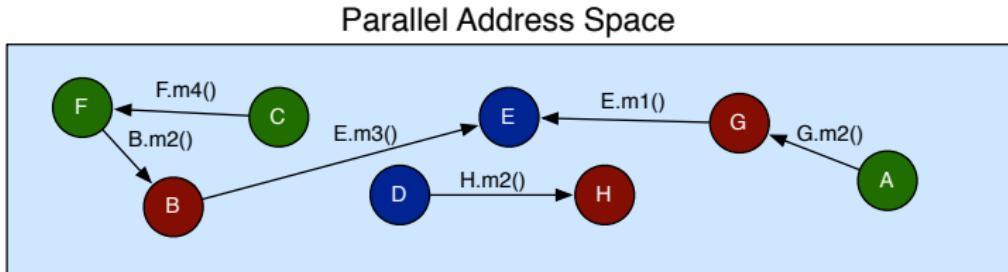
Benefits

Charm++

Primer

Collections

Conclusion



- ▶ How can objects communicate across address spaces?
  - ▶ Just like a sequential object-oriented language, an object's reference is used to invoke a method
  - ▶ In the parallel space, this is a handle that is location transparent
  - ▶ A method invocation becomes an act of communication

# Method-driven Asynchronous Communication

## Message-driven execution

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

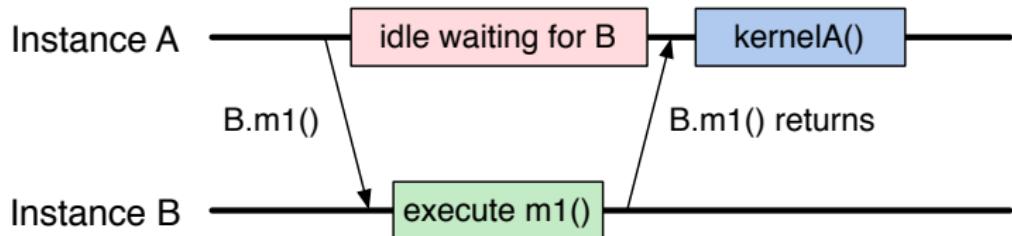
Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion



- ▶ What happens if an object waits for a return value from a method invocation?
  - ▶ Performance
  - ▶ Latency
  - ▶ Reasoning about correctness

# Design Principle

Do not wait for remote completion

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

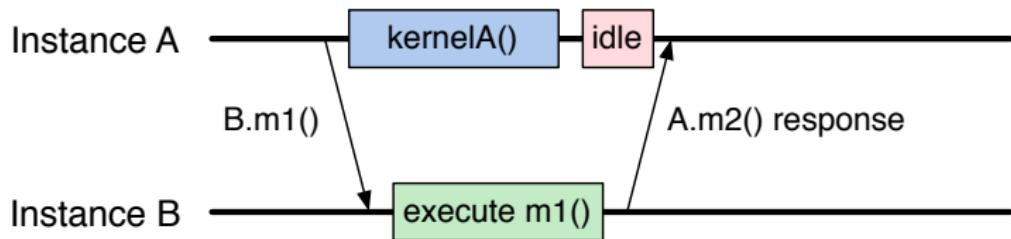
Object Design  
Execution Mode

Hello World

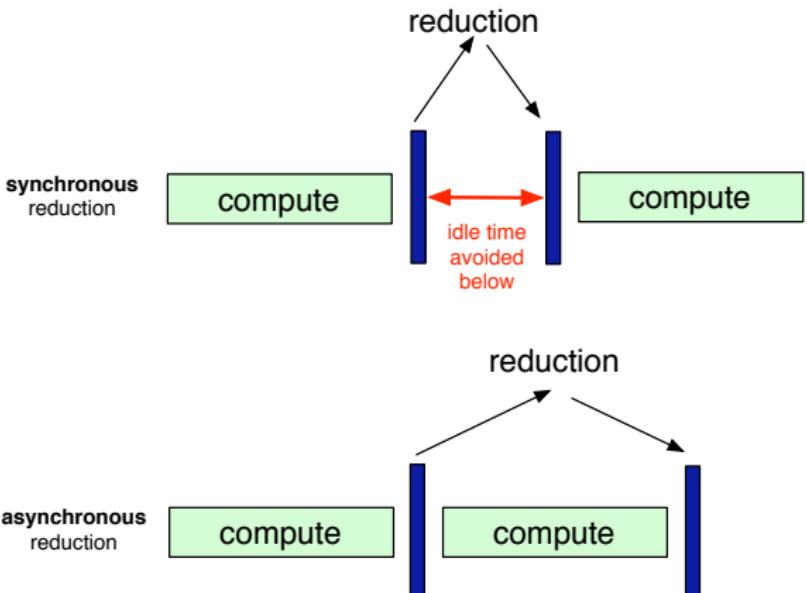
Benefits

Charm++  
Primer  
Collections

Conclusion



- ▶ Hence, method invocations should be asynchronous
  - ▶ No return values
- ▶ Computations are driven by the incoming data
  - ▶ Initiated by the sender or method caller



# Methods

Natural units of sequential computation

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

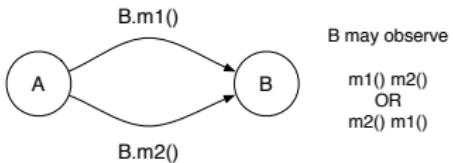
Hello World

Benefits

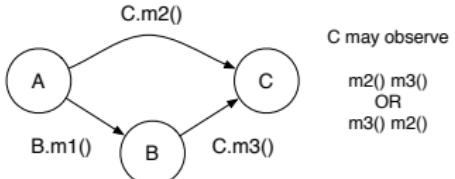
Charm++  
Primer  
Collections

Conclusion

- ▶ Methods still have the same sequential semantics
  - ▶ Atomicity: methods do not execute in parallel
- ▶ Methods cannot be interrupted or preempted
- ▶ Methods interact and update state of an object in the same way
- ▶ Method sequencing is what changes from sequential computation



B may observe  
 $m1() \ m2()$   
OR  
 $m2() \ m1()$

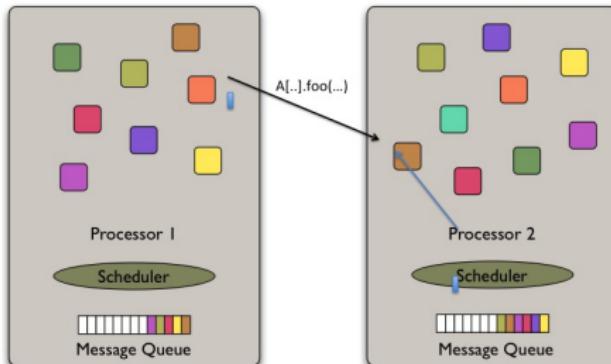


C may observe  
 $m2() \ m3()$   
OR  
 $m3() \ m2()$

# The Execution Model

## Messages and schedulers

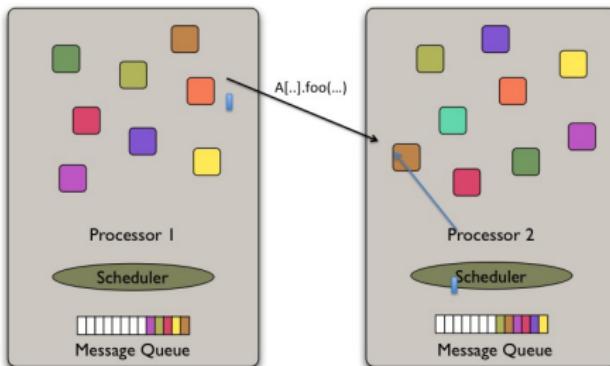
- ▶ Several objects live on a single *processing entity (PE)*
  - ▶ For now, think of it as a core (or just “processor”)
- ▶ As a result,
  - ▶ Method invocations directed at objects on that processor will have to be stored in a pool
  - ▶ A user-level scheduler will select one invocation from the queue and runs it to completion
  - ▶ A PE is the entity that has one scheduler instance associated with it



# Message-driven Execution

Objects are reactive entities

- ▶ Execution is triggered by availability of a “message” (a method invocation)
- ▶ When an entry method executes:
  - ▶ it may generate messages for other objects
  - ▶ the *runtime system* (RTS) deposits them in the message Q on the target processor



# Example

## Hello World

- ▶ **hello.ci** file

```
mainmodule hello {
    mainchare Main {
        entry Main(CkArgMsg *m);
    };
}
```

- ▶ **hello.cpp** file

```
#include <stdio.h>
#include "hello.decl.h"

class Main : public CBase_Main {
public: Main(CkArgMsg* m) {
    ckout << "Hello World!" << endl;
    CkExit();
}
};

#include "hello.def.h"
```

# Example

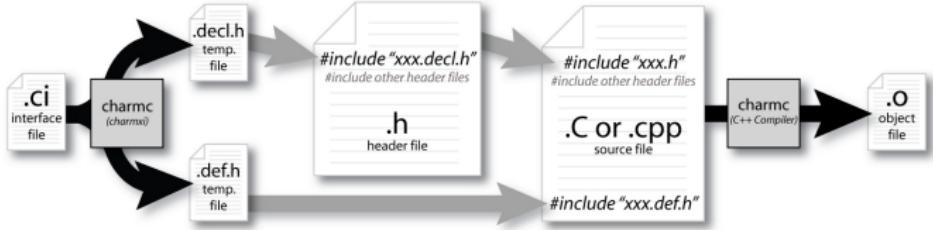
## Hello World with chares

hello.ci file

```
mainmodule hello {  
    mainchare Main {  
        entry Main(CkArgMsg *m  
        );  
    };  
    chare Singleton {  
        entry Singleton();  
    };  
};
```

hello.cpp file

```
#include <stdio.h>  
#include "hello.decl.h"  
  
class Main : public CBase_Main {  
public: Main(CkArgMsg* m) {  
    CProxy_Singleton::ckNew();  
};  
};  
  
class Singleton : public  
    CBase_Singleton {  
public: Singleton() {  
    ckout << "Hello World!" <<  
        endl;  
    CkExit();  
};  
};  
#include "hello.def.h"
```



# Compilation Example

## Hello World

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Compiling
  - ▶ `charmcc hello.ci`
  - ▶ `charmcc -c hello.C`
  - ▶ `charmcc -o hello hello.o`
- ▶ Running
  - ▶ `./charmrun +p7 ./hello`
  - ▶ The `+p7` tells the system to use seven cores

# Building Charm++

## 1. Download the latest stable version of the code:

```
http://charmplusplus.org/download/
```

## 2. If you are feeling adventurous, get the latest version:

```
git clone http://charm.cs.uiuc.edu/gerrit/charm
```

## 3. Build an appropriate version of the code for your architecture:

```
./build <TARGET> <ARCH> <OPTS>
```

- ▶ TARGET = charm++, AMPI, bgmpi, LIBS, etc.
- ▶ ARCH = net-linux-x86\_64, multicore-darwin-x86\_64, pamilrts-bluegeneq, etc.
- ▶ OPTS = --with-production, --enable-tracing, xlc, smp, -j8, etc.

## 4. Get familiar with Charm++:

```
http://charmplusplus.org/tutorial/
```

## 5. Ask questions:

```
http://charm.cs.illinois.edu/manuals/html/charm++/A.html
```

# Exercise 1.1

## Building Charm++

Download Charm++ code base and build it for your target machine (laptop, desktop, server, supercomputer)

### Examples:

- ▶ Linux cluster:

```
./build charm++ netlrts-linux-x86_64 gcc -j8 --with-production
```

- ▶ Linux cluster on MPI:

```
./build charm++ mpi-linux-x86_64 -j8 --with-production
```

- ▶ Linux laptop:

```
./build charm++ multicore-linux-x86_64 -j8 --with-production
```

- ▶ Mac laptop:

```
./build charm++ multicore-darwin-x86_64 -j8 --with-production
```

- ▶ Windows laptop with VC++:

```
./build charm++ multicore-win-x86_64 --with-production
```

# Exercise 1.2

## Running Charm++ Hello World

Get into Charm++'s source code and look for the Hello World example:

`charm-6.8.0/examples/charm++/hello`

Run the example

# Impact on Communication

## Alleviating one performance bottleneck

- ▶ Current use of communication network:
  - ▶ Compute-communicate cycles in typical MPI apps
  - ▶ Network is used for a fraction of time
  - ▶ On the critical path
- ▶ Hence, current communication networks are over-engineered by necessity
- ▶ With overdecomposition:
  - ▶ Communication is spread over an iteration
  - ▶ Adaptive overlap of communication and computation

# Example

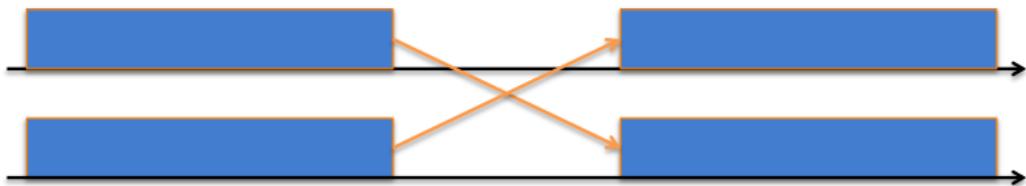
## Stencil computation

- ▶ Consider a simple stencil computation
  - ▶ With traditional design based on traditional methods (e.g. MPI-based)
    - ▶ Each processor has a chunk, which alternates between computing and communicating
  - ▶ With Charm++
    - ▶ Multiple chunks on each processor
    - ▶ Wait time for each chunk overlapped with useful computation for others
    - ▶ Communication spread over time

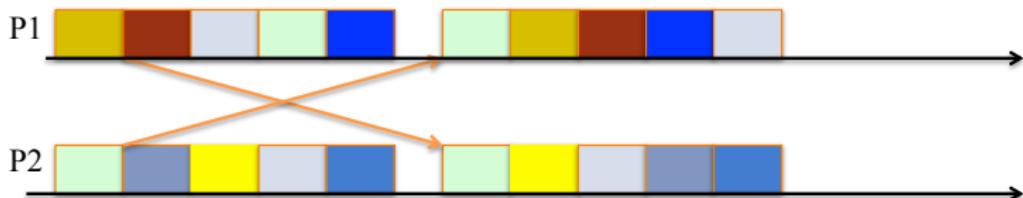
# Example

## Stencil computation

Stencil in MPI: No overlap among computation and communication

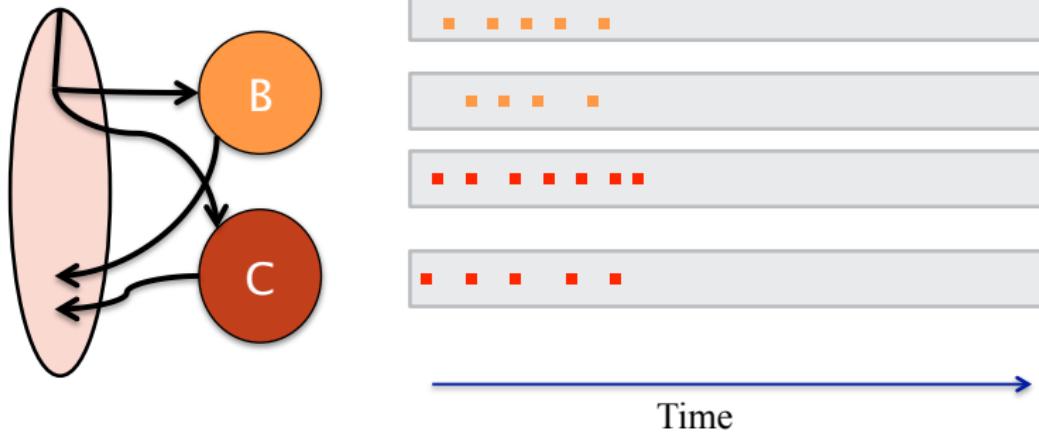


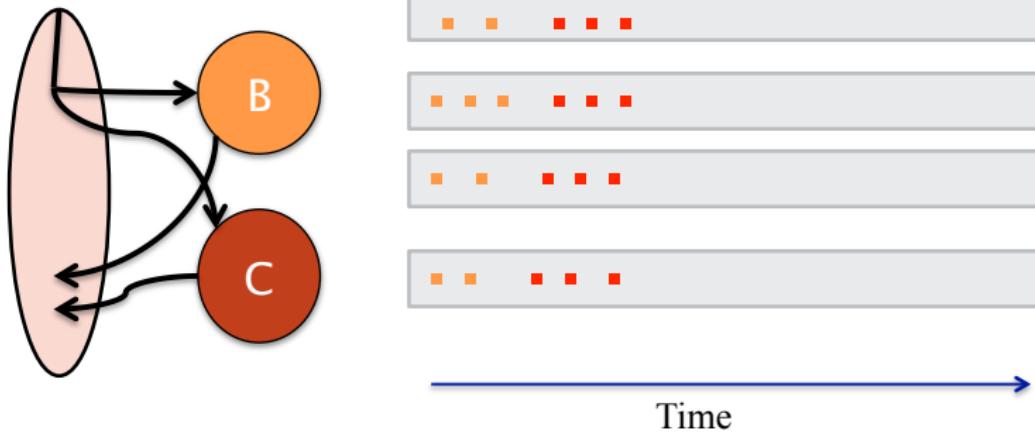
Stencil in Charm: Communication of a chare overlaps with computation of others



# Space Division

An effect of lacking message-driven execution (and virtualization)

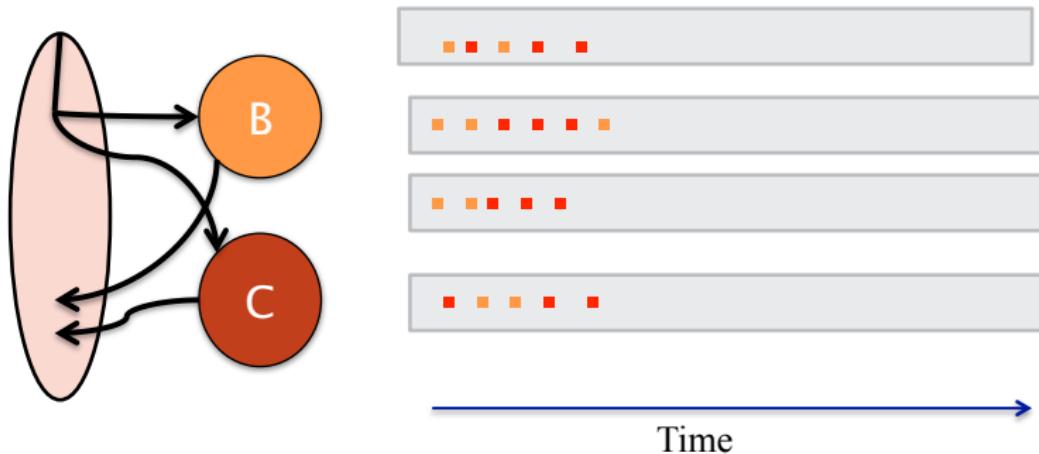




# Modularity and Composability

Separation of concerns

Parallel Composition: A1; (B || C ); A2



Different modules, written in different languages/paradigms, can overlap in time and on processors, without programmer having to worry about this explicitly

# Migratability

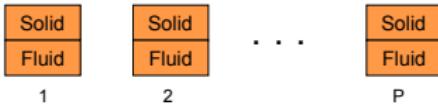
## Location independence

- ▶ Once the programmer has written the code without reference to processors, all of the communication is expressed among objects
- ▶ The system is free to migrate the objects across processors as and when it pleases
  - ▶ It must ensure it can deliver method invocations to the objects, wherever they go
  - ▶ This migratability turns out to be a key attribute for empowering an adaptive runtime system

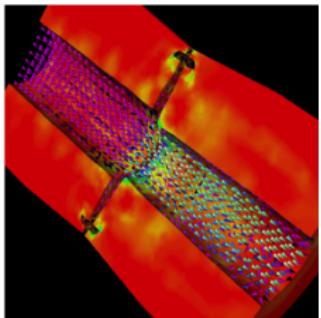
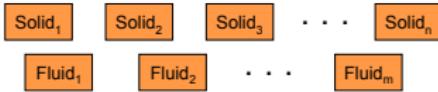
# Example

Decomposition independent of number of cores

- ▶ Rocket simulation under traditional MPI



- ▶ Rocket simulation with migratable objects



- ▶ Benefits: load balance, communication optimizations, modularity

# Architectures

Utility for multi-cores, many-cores, accelerators

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

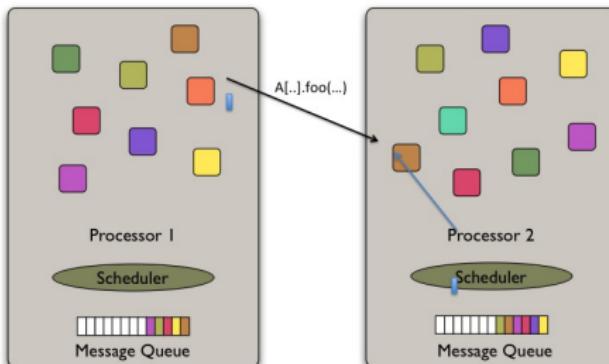
Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion



# Load Balancing

A major benefit of migratability

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

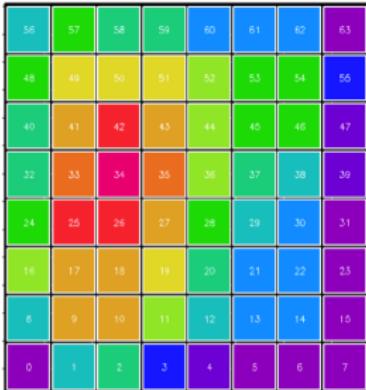
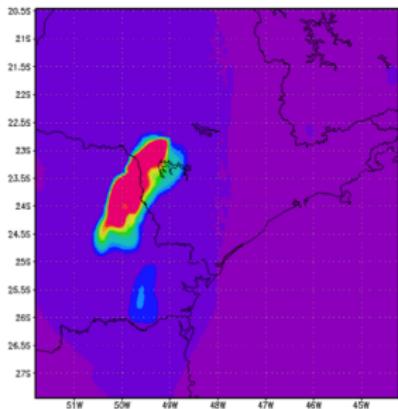
- ▶ Static
  - ▶ Irregular applications
  - ▶ Programmer shouldn't have to figure out ideal mapping
- ▶ Dynamic
  - ▶ Applications are increasingly using adaptive strategies
  - ▶ Abrupt refinements
  - ▶ Continuous migration of work (e.g. particles in MD)
- ▶ Challenges
  - ▶ Performance limited by most overloaded processor
  - ▶ The chance that one processor is severely overloaded gets higher as processor count increases

**Migratable objects empower automated load balancing!**

# Example

## Weather forecasting in BRAMS

- ▶ BRAMS: Brazilian weather code (based on RAMS)
- ▶ AMPI version (Eduardo Rodrigues, with Dr. Celso Mendes and Prof. Jairo Panetta)



# Basic Virtualization of BRAMS

Data partition

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

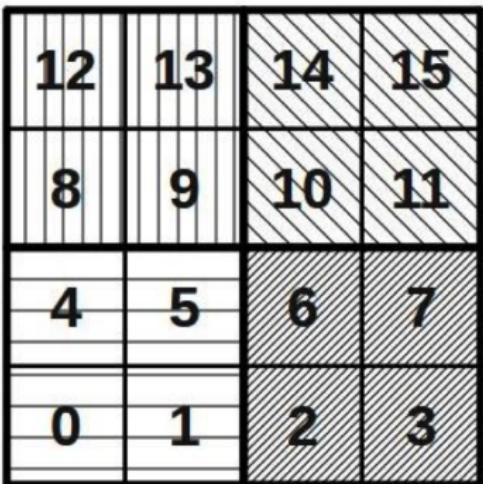
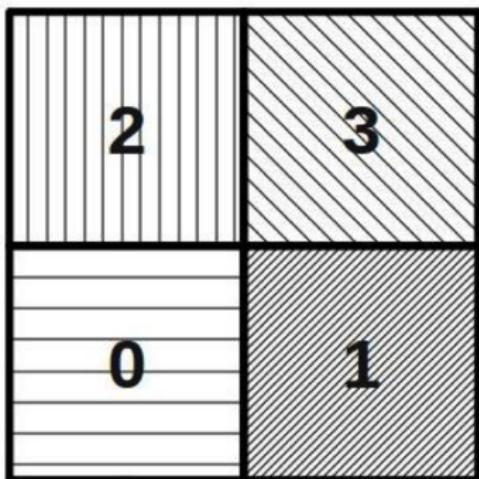
Object Design  
Execution Mode

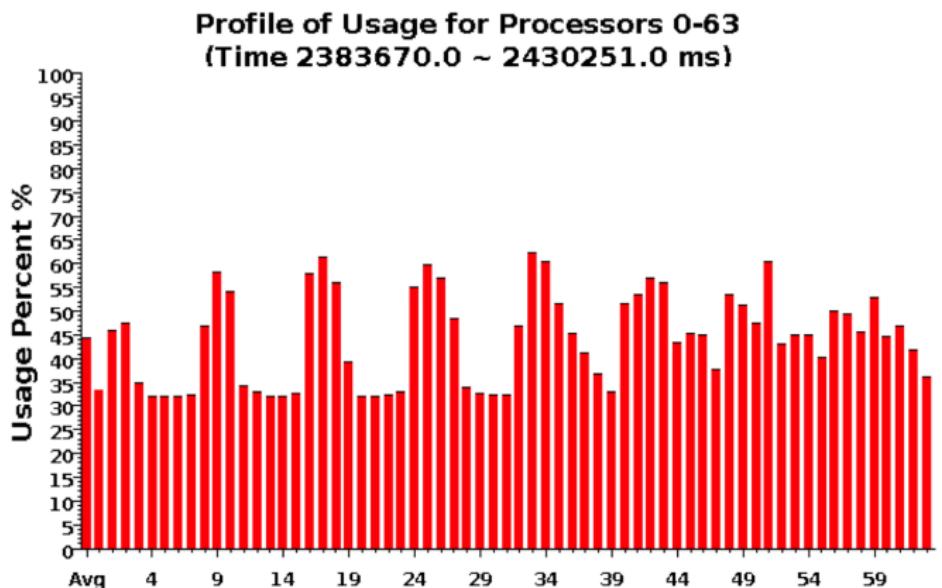
Hello World

Benefits

Charm++  
Primer  
Collections

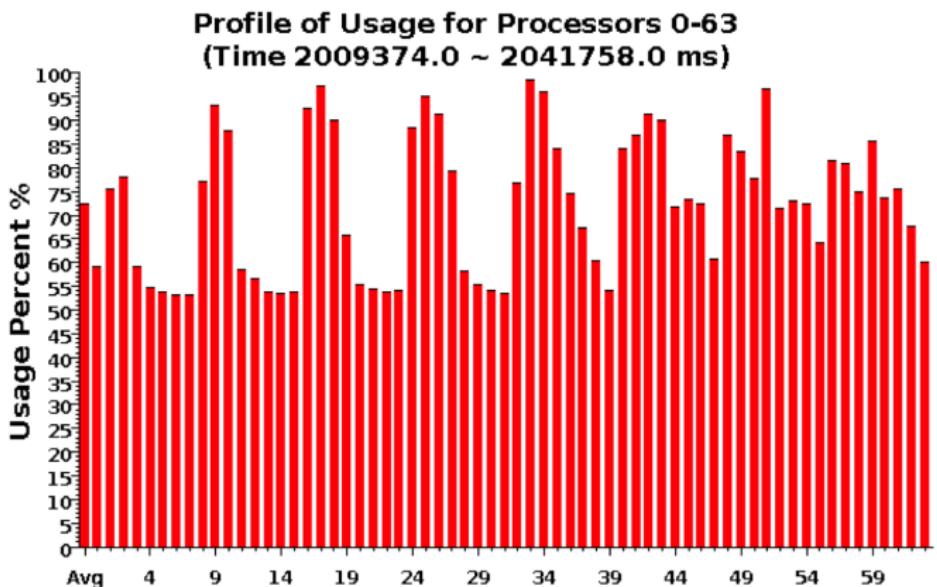
Conclusion





# Over-decomposition

1024 objects on 64 processors

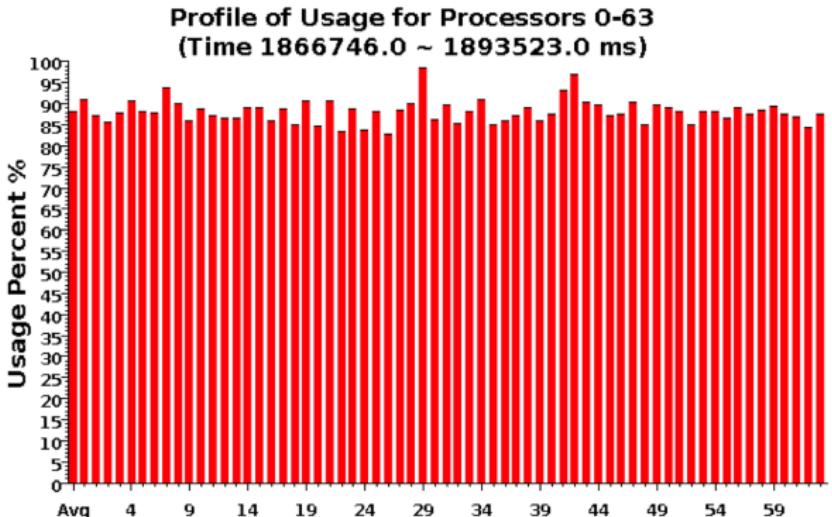


Benefits from communication/computation overlap

# With Load Balancing

1024 objects on 64 processors

- ▶ No overdecomposition (64 threads): 4988 sec
- ▶ Overdecomposition into 1024 threads: 3713 sec
- ▶ Load balancing (1024 threads): 3367 sec



# Summary of Benefits

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

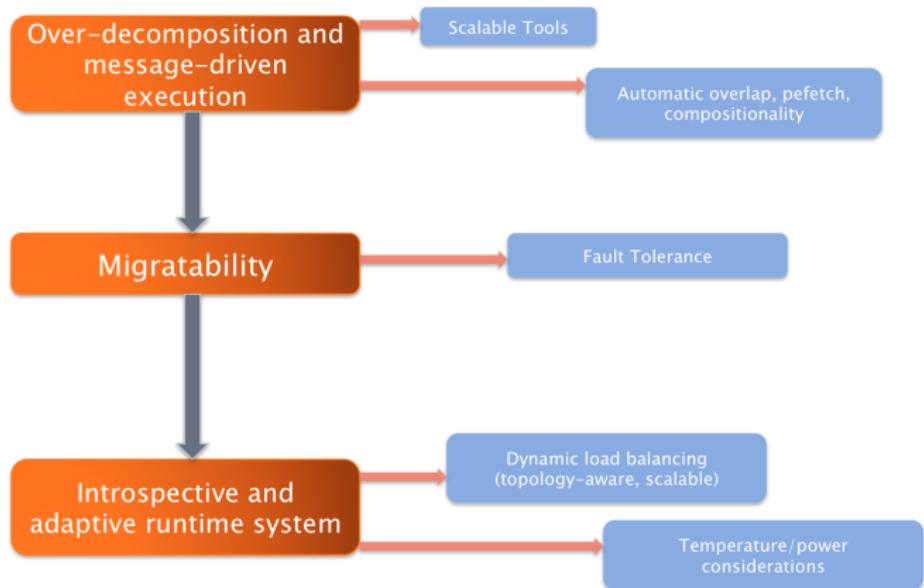
Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion



# The Charm++ Programming Language

## History

University of Illinois at Urbana-Champaign, early 1990s



Prof. Laxmikant V. Kalé



Prof. Klaus Schulten



Parallel Programming Laboratory (Kalé)

Theoretical and Computational Biophysics Group (Schulten)

# Parallel Programming Galore

## List of Parallel Programming Languages in the 90s

"C* in C	CUMULVS	Java RMI	P-RIO	Quake
ABCPL	DAGGER	javaPG	P3L	Quark
ACE	DAPPLE	JAVAR	P4-Linda	Quick Threads
ACT++	Data Parallel C	JavaSpaces	Pablo	Sage++
ADDAP	DC++	JIDL	PADE	SAM
Adl	DCE++	Joyce	PADRE	SCANDAL
Adsmith	DDD	Karma	Panda	SCHEDULE
AFAPI	DICE	Khoros	Papers	SciTL
ALWAN	DIFC	KOAN/Fortran-S	Para++	SDDA
AM	Distributed Smalltalk	LAM	Paradigm	SHMEM
AMDC	DOLIB	Legion	Parafuse2	SIMPLE
Amoebs	DOME	Lilac	Paradise	Sina
AppleS	DOSMOS	Linda	Paralaxis	SISAL
ARTS	DRL	LIPS	Parallel Haskell	SMI
Athropuscan-Ob	DSM-Threads	Locust	Parallel-C++	SONIC
Aurora	Ease	Lparx	ParC	Split-C
Autonmap	ECO	Lucid	ParLib++	SR
bb_THREADS	Eilean	Maisie	ParLin	SThreads
Blaze	Emerald	Manifold	Parlog	Strand
BlockComm	EPL	Mentat	Parmacs	SUFP
BSP	Excalibur	Meta Chacos	Parti	SuperPascal
C*	Express	Midway	pC	Synergy
C**	Falcon	Millipede	pC++	TCGMSG
C4	Filaments	Mirage	PCN	Telegraphos
Carlos	FLASH	Modula-2*	PCP	The FORCE
Cashmere	FM	Modula-P	PCU	Threads.h++
CC++	Fork	MOSIX	PEACE	TRAPPER
Charlotte	Fortran-M	MpC	PENNY	TreadMarks
Charm	FX	MPC++	PET	UC
Charm++	GAMMA	MPI	PETSc	uC++
Chi	GLAMMA	Multipol	PI	UNITY
Cilk	Glenda	Mura	Phosphorus	V
CM-Fortran	GUARD	Nano-Threads	POET	Vice
Code	HAIL	NEST	POLE	Virtual V-NUS
Concurrent ML	HORUS	NetClasses++	POOL-T	VPE
Converse	HPC	Noe	POOMA	Win32 threads
COOL	HPC++	Nimrod	POSYBL	WinPar
CORRELATE	HPP	NOW	PRESTO	WWWindz
CpaPar	IMPACT	Objective Linda	Prospero	XENOOPS
CPI*	ISETL-Linda	Occam	Proteus	XPC
CRL	OSS	Omega	PSDM	Zounds
CSP	JADA	OOF90	PSI	ZPI
Cthreads	JADE	P++	PVM	
			QPC++	

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++

Primer  
Collections

Conclusion

"Patterns for Parallel Programming" (Tim Mattson, et al, 2004)

# Charm++ Runtime System

A whole programming ecosystem

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

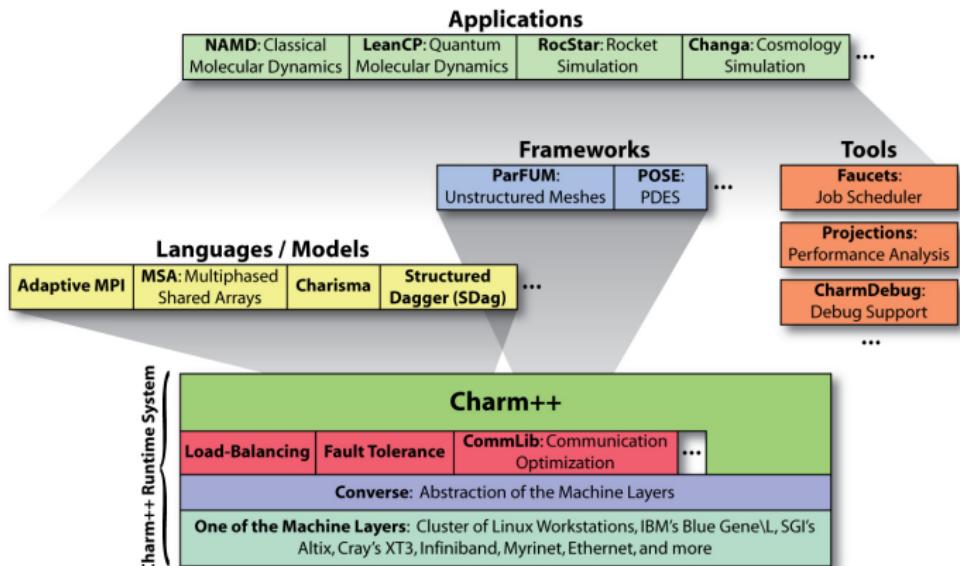
Hello World

Benefits

Charm++

Primer  
Collections

Conclusion



# Charm++ File Structure

Header, source, and interface files

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

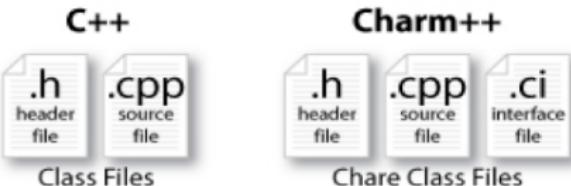
Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ C++ objects (including Charm++ objects)
  - ▶ Defined in regular .h and .C files
- ▶ Chare objects, entry methods (asynchronous methods)
  - ▶ Defined in .ci file
  - ▶ Implemented in the .C file



# Charm Interface

## Modules

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Charm++ programs are organized as a collection of modules
- ▶ Each module has one or more chares
- ▶ The module that contains the *mainchare*, is declared as the `mainmodule`
- ▶ Each module, when compiled, generates two files:  
`MyModule.decl.h` and `MyModule.def.h`

.ci file

```
[main]module MyModule {  
    //... chare definitions ...  
};
```

# Charm Interface

## Chares

- ▶ Chares are parallel objects that are managed by the RTS
- ▶ Each chare has a set *entry methods*, which are asynchronous methods that may be invoked remotely
- ▶ The following code, when compiled, generates a C++ class `CBase_MyChare` that encapsulates the RTS object
- ▶ This generated class is extended and implemented in the .C file

### .ci file

```
[main]chare MyChare {  
    //... entry method definitions ...  
};
```

### .C file

```
class MyChare : public CBase_MyChare {  
    //... entry method implementations ...  
};
```

# Charm Interface

## Entry methods

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++

Primer  
Collections

Conclusion

- ▶ Entry methods are C++ methods that can be remotely and asynchronously invoked by another chare

.ci file:

```
entry MyChare(); /* constructor entry method */  
entry void foo();  
entry void bar(int param);
```

.C file:

```
MyChare::MyChare() { /*... constructor code ...*/ }  
  
MyChare::foo() { /*... code to execute ...*/ }  
  
MyChare::bar(int param) { /*... code to execute ...*/ }
```

# Charm Interface

`mainchare`

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Execution begins with the `mainchare`'s constructor
- ▶ The `mainchare`'s constructor takes a pointer to system-defined class `CkArgMsg`
- ▶ `CkArgMsg` contains `argv` and `argc`
- ▶ The `mainchare` will typically create some additional chares

# Creating a Chare

Interacting object on a distributed system

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++

Primer  
Collections

Conclusion

- ▶ A chare declared as `chare MyChare { ... };` can be instantiated by the following call:

```
CProxy_MyChare::ckNew(... constructor arguments ...);
```

- ▶ To communicate with this class in the future, a *proxy* to it must be retained

```
CProxy_MyChare proxy =  
CProxy_MyChare::ckNew(... constructor arguments ...);
```

# Chare Proxies

Representative of a real object

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++

Primer  
Collections

Conclusion

- ▶ A chare's own proxy can be obtained through a special variable `thisProxy`
- ▶ Chare proxies can also be passed so chares can learn about others
- ▶ In this snippet, `MyChare` learns about a chare instance `main`, and then invokes a method on it:

.ci file

```
entry void foobar2(CProxy_Main main);
```

.C file

```
MyChare::foobar2(CProxy_Main main) {
    main.foo();
}
```

# Charm++ Termination

More than finishing execution

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ There is a special system call `CkExit()` that terminates the parallel execution on all processors (but it is called on one processor) and performs the required cleanup
- ▶ The traditional `exit()` is insufficient because it only terminates one process, not the entire parallel job (and will cause a hang)
- ▶ `CkExit()` should be called when you can safely terminate the application (you may want to synchronize before calling this)

```
mainmodule MyModule {
    mainchare Main {
        entry Main(CkArgMsg *m);
    };

    chare Simple {
        entry Simple(int x, double y);
    };
}
```

# Example

## Chare creation, .C file

```
#include <stdio.h>
#include "MyModule.decl.h"

class Main : public CBase_Main {
public: Main(CkArgMsg* m) {
    ckout << "Hello World!" << endl;
    if (m->argc > 1) ckout << " Hello " << m->argv[1] << "!!!"
        << endl;
    double pi = 3.1415;
    CProxy_Simple::ckNew(12, pi);
}
};

class Simple : public CBase_Simple {
public: Simple(int x, double y) {
    ckout << "Hello from a simple chare running on " << CkMyPe()
        << endl;
    ckout << "Area of a circle of radius" << x << " is " << y*x*x
        << endl;
    CkExit();
}
};

#include "MyModule.def.h"
```

# Asynchronous Methods

Beyond two-sided communication

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++

Primer  
Collections

Conclusion

- ▶ Entry methods are invoked by performing a C++ method call on a chare's proxy

```
CProxy<MyChare> proxy =  
    CProxy<MyChare>::ckNew(... constructor arguments ...);  
  
proxy.foo();  
proxy.bar(5);
```

- ▶ The `foo` and `bar` methods will then be executed with the arguments, wherever the created chare, `MyChare`, happens to live
- ▶ The policy is one-at-a-time scheduling (that is, one entry method on one chare executes on a processor at a time)

# Asynchronous Methods

## Beyond FIFO channels

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Method invocation is not ordered (between chares, entry methods on one chare, etc.)
- ▶ For example, if a chare executes this code:

```
CProxy_MyChare proxy = CProxy_MyChare::ckNew();
proxy.foo();
proxy.bar(5);
```

- ▶ These prints may occur in **any** order

```
MyChare::foo() {
    ckout << "foo executes" << endl;
}

MyChare::bar(int param) {
    ckout << "bar executes with " << param << endl;
}
```

# Example

## Asynchronous methods

- ▶ For example, if a chare invokes the same entry method twice:

```
proxy.bar(7);  
proxy.bar(5);
```

- ▶ These may be delivered in **any** order

```
MyChare::bar(int param) {  
    ckout << "bar executes with " << param << endl;  
}
```

- ▶ Output

```
bar executes with 5  
bar executes with 7
```

OR

```
bar executes with 7  
bar executes with 5
```

```
mainmodule MyModule {
    mainchare Main {
        entry Main(CkArgMsg *m);
    };
    chare Simple {
        entry Simple(double y);
        entry void findArea(int radius, bool done);
    };
}
```

# Exercise 1.3

Does this program execute correctly?

```
struct Main : public CBase_Main {
    Main(CkArgMsg* m) {
        double pi = 3.1415;
        CProxy_Simple sim = CProxy_Simple::ckNew(pi);
        for (int i = 1; i< 10; i++) sim.findArea(i, false);
        sim.findArea(10, true);
    };
};

struct Simple : public CBase_Simple {
    float y;
    Simple(double pi) {
        y = pi;
        ckout << "Hello from a simple chare running on " << CkMyPe()
            << endl;
    }
    void findArea(int r, bool done) {
        ckout << "Area of a circle of radius" << r << " is " << y*r*r <<
            endl;
        if (done) CkExit();
    }
};
```

# Data Types and Entry Methods

STL supported

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++

Primer  
Collections

Conclusion

- ▶ You can pass basic C++ types to entry methods (`int`, `char`, `bool`, etc.)
- ▶ C++ STL data structures can be passed by including `pup_stl.h`
- ▶ Arrays of basic data types can also be passed like this:
- ▶ .ci file:

```
entry void foobar(int length, int data[length]);
```

- ▶ .C file:

```
MyChare::foobar(int length, int* data) {  
    // ... foobar code ...  
}
```

# Readonlys

## Constants on a distributed system

- ▶ A *readonly* is a global (within a module) read-only variable that can only be written to in the `mainchare`'s constructor
- ▶ Can then be read (**not written!**) by any chare in the module
- ▶ It is declared in the .ci file:

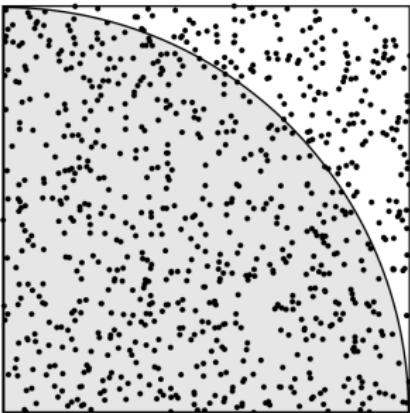
```
readonly <type> <name>;
readonly CProxy_Main mainProxy;
readonly int numChares;
```

- ▶ And defined in the .C file:

```
<type> <name>;
CProxy_Main mainProxy;
int numChares;
```

- ▶ And set in the `mainchare`'s constructor

```
MyChare::MyChare(CkArgMsg *m) {
    mainProxy = thisProxy;
    numChares = 10;
}
```



Complete the code in directory:

code/pi/skeleton

of the class repository:

```
git clone https://github.com/emenesesrojas/parallel_objects.git
```

Do not forget to define variable CHARMDIR in the Makefile

# Collections of Objects

## Concepts

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Objects can be grouped into indexed collections
- ▶ Basic examples
  - ▶ Matrix block
  - ▶ Chunk of unstructured mesh
  - ▶ Portion of distributed data structure
  - ▶ Volume of simulation space
- ▶ Advanced examples
  - ▶ Abstract portions of computation
  - ▶ Interactions among basic objects or underlying entities

# Collections of Objects

## Types of collections

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Structured: 1D, 2D, ..., 6D
- ▶ Unstructured: anything hashable
- ▶ Dense
- ▶ Sparse
- ▶ Static - all created at once
- ▶ Dynamic - elements come and go

```
mainmodule arr {  
  
    mainchare Main {  
        entry Main(CkArgMsg*);  
    }  
  
    array [1D] hello {  
        entry hello(int);  
        entry void printHello();  
    }  
}
```

# Example

## Chare array Hello World

```
#include "arr.decl.h"

struct Main : CBase_Main {
    Main(CkArgMsg* msg) {
        int arraySize = atoi(msg->argv[1]);
        CProxy_hello p = CProxy_hello::ckNew(arraySize, arraySize);
        p[0].printHello();
    }
};

struct hello : CBase_hello {
    hello(int n) : arraySize(n) { }
    hello(CkMigrateMessage*) { }
    void printHello() {
        CkPrintf("PE[%d]: hello from p[%d]\n", CkMyPe(), thisIndex);
        if (thisIndex == arraySize - 1) CkExit();
        else thisProxy[thisIndex + 1].printHello();
    }
    private:
    int arraySize;
};

#include "arr.def.h"
```

# Example

## Hello World array Projections timeline view



Run on BG/Q 16 Nodes, mode c16, 1024 elements (4 per process)

# Declaring a Chare Array

## Dimensionality and functionality

.ci file:

```
array [1d] foo {  
    entry foo(); // constructor  
    // ... entry methods ...  
}  
array [2d] bar {  
    entry bar(); // constructor  
    // ... entry methods ...  
}
```

.C file:

```
struct foo : public CBase_foo {  
    foo() { }  
    foo(CkMigrateMessage*) { }  
    // ... entry methods ...  
};  
struct bar : public CBase_bar {  
    bar() { }  
    bar(CkMigrateMessage*) { }  
    // ... entry methods ...  
};
```

# Constructing a Chare Array

Elements are distributed on nodes of the system

- ▶ Constructed much like a regular chare
- ▶ The size of each dimension is passed to the constructor

```
void someMethod() {  
    CProxy_foo::ckNew(10);  
    CProxy_bar::ckNew(5, 5);  
}
```

- ▶ The proxy may be retained:

```
CProxy_foo myFoo = CProxy_foo::ckNew(10);
```

- ▶ The proxy represents the entire array, and may be indexed to obtain a proxy to an individual element in the array

```
myFoo[4].invokeEntry();
```

# thisIndex

Pretty much like this

- ▶ 1d: `thisIndex` returns the index of the current char array element
- ▶ 2d: `thisIndex.x` and `thisIndex.y` returns the indices of the current char array element

.ci file:

```
array [1d] foo {
    entry foo();
}
```

.C file:

```
struct foo : public CBase_foo {
    foo() {
        CkPrintf("array index = %d", thisIndex);
    }
};
```

# Collections of Objects

Runtime service

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ System knows how to ‘find’ objects efficiently:  
 $(collection, index) \rightarrow processor$
- ▶ Applications can specify a mapping, or use simple runtime-provided options (e.g. blocked, round-robin)
- ▶ Distribution can be static, or dynamic
- ▶ Key abstraction: application logic doesn’t change, even though performance might

# Collections of Objects

Runtime service

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

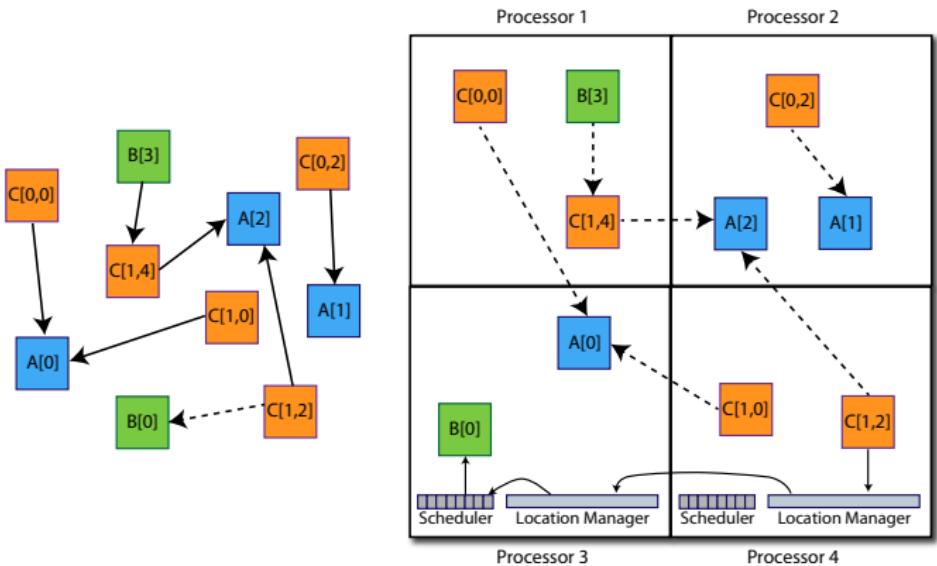
Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Can develop and test logic in objects separately from their distribution
- ▶ Separation in time: make it work, then make it fast
- ▶ Division of labor: domain specialist writes object code, computationalist writes mapping
- ▶ Portability: different mappings for different systems, scales, or configurations
- ▶ Shared progress: improved mapping techniques can benefit existing code



# Collective Communication Operations

Efficiency and expressivity

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

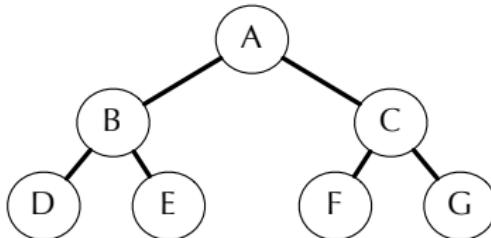
Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion



# Broadcast

## One-to-many

- ▶ A message to each object in a collection
- ▶ The chare array proxy object is used to perform a broadcast
- ▶ It looks like a function call to the proxy object
- ▶ From the main chare:

```
CProxy>Hello helloArray = CProxy>Hello::ckNew(helloArraySize);  
helloArray.foo();
```

- ▶ From a chare array element that is a member of the same array:

```
thisProxy.foo()
```

- ▶ From any chare that has a proxy p to the chare array

```
p.foo()
```

# Reduction

## Many-to-one

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

- ▶ Combines a set of values: sum, max, aggregate
- ▶ Usually reduces the set of values to a single value
- ▶ Combination of values requires an operator
- ▶ The operator must be commutative and associative
- ▶ Each object calls `contribute` in a reduction

# Example

## Reduction

Programming with  
Parallel Objects

Esteban Meneses

HPC Challenges

Concepts

Object Design  
Execution Mode

Hello World

Benefits

Charm++  
Primer  
Collections

Conclusion

```
mainmodule reduction {
    mainchare Main {
        entry Main(CkArgMsg* msg);
        entry [reductiontarget] void done(int value);
    };
    array [1D] Elem {
        entry Elem(CProxy_Main mProxy);
    };
}
```

# Example

## Reduction

```
#include "reduction.decl.h"

const int numElements = 49;

class Main : public CBase_Main {
public:
    Main(CkArgMsg* msg) { CProxy_Elem::ckNew(thisProxy, numElements); }
    void done(int value) {
        CkAssert(value == numElements * (numElements - 1) / 2);
        CkPrintf("value: %d\n", value);
        CkExit();
    }
};

class Elem : public CBase_Elem {
public:
    Elem(CProxy_Main mProxy) {
        int val = thisIndex;
        CkCallback cb(CkReductionTarget(Main, done), mProxy);
        contribute(sizeof(int), &val, CkReduction::sum_int, cb);
    }
    Elem(CkMigrateMessage*) { }
};

#include "reduction.def.h"
```

## Output:

```
value: 1176
Program finished.
```

# Exercise 1.5

## Key-value store

Implement a key-value store using a 1D char array:

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

Complete the code in directory:

code/keyvalue/skeleton

of the class repository:

```
git clone https://github.com/emenesesrojas/parallel_objects.git
```

Do not forget to define variable CHARMDIR in the Makefile

## Concluding Remarks

- ▶ Parallel objects bring about a modular and composable parallel programming framework
- ▶ Two major benefits of over-decomposition are:
  - ▶ Overlap of communication and computation
  - ▶ Load balancing
- ▶ Chare arrays are the main object collection in Charm++

**Thank You!**  
**Q&A**

